

Lecture 10 boosting

wbg231

December 2022

1 motivation

- recall in ada boosting we learn weak learners $g_m(x)$ and weights α_m and use those to get our final predictor $G(x) = \text{sing}(\sum_{i=1}^m \alpha_m G_m(x))$
- why not just learn $G(x)$ directly?
- well it is a linear sum of weak learners that are not necessarily linear so we learn it sequentially in ada boost

nonlinear regression

- we can find wide types of data by fitting a linear combination of transformations to the input

$$f(x) = \sum_{m=1}^m v_m h_m(x)$$

where h_m is called a basis function such that

$$h_1 \cdots h_m : X \rightarrow \mathbb{R}$$

- example polynomial regression $h_i(x) \in \{x^i : i \in \mathbb{Z}\}$
- we can fit this with standard linear models if our basis functions are fixed ahead of time all we are learning are weighting factors

adaptive bayes function model

- what if we want to learn the basis functions
- we define a base hypotheses space $\mathcal{H} : X \rightarrow \mathbb{R}$ so all scalar values function

- an [adaptive basis function expansion](#) over \mathcal{H} is an ensemble model

$$f(x) = \sum_m^M v_m h_m(x)$$

where $v_m \in \mathbb{R}, h_m \in \mathcal{H}$

- so then we can combine these to get a new hypotheses space

$$\mathcal{F}_M = \left\{ \sum_m v_m h_m(x) \mid v_m \in \mathbb{R}, h_m \in \mathcal{H} \forall m \in [1 \dots M] \right\}$$

- so we our objective is

$$j(v_1 \dots v_m, h_1 \dots h_m) = \frac{1}{n} \ell(y_i, f(x_i)) = \frac{1}{n} \ell(y_i, \sum_{m=1}^M v_m h_m(x))$$

- if we want to optimize this sometimes we can use gradient descent or find a closed form (but that is not always the case)
- in cases where we can not differentiate we can try using a greedy algorithm similar to ada boost

gradient boosting

- applies when ever our loss function is sub differentiable wrt our training predictions $f(x_i)$ we can do regressions with the hypothesis base space \mathcal{H}

forward Stagewise adaptive modeling

- to recap our goal is to find the model

$$f(x) = \sum_m v_m h_m(x)$$

that is a weighed sum of basis functions given some loss function

- we do this by greedily fitting one function at a time without adjusting previous functions "forward Stagewise"
- so after $m - 1$ stages we will have

$$f_{m-1} = \sum_{i=1}^{m-1} v_i h_i$$

- and then at the m th round we are trying to find the basis function $h \in \mathcal{H}$ and $v_m > 0$ such that

$$f_m = f_{m-1} + v_m h_m = c + v_m h_m$$

improves our loss as much as possible

- so this is what our algorithm looks like

let's plug in our objective function.

1 Initialize $f_0(x) = 0$.

2 For $m = 1$ to M :

1 Compute:

$$(v_m, h_m) = \arg \min_{v \in \mathbb{R}, h \in \mathcal{H}} \frac{1}{n} \sum_{i=1}^n \ell \left(y_i, f_{m-1}(x_i) + \underbrace{v h(x_i)}_{\text{new piece}} \right).$$

2 Set $f_m = f_{m-1} + v_m h_m$.

3 Return: f_M .

- we are going to set our loss function as exponential ie $\ell(y, f(x)) = e^{-yf(x)}$
- and assume our $\mathcal{H} = \{h : x \rightarrow \{-1, 1\}\}$ that is our base functions are binary classifiers

$$\begin{aligned} J(v, h) &= \sum_{i=1}^n \exp[-y_i (f_{m-1}(x_i) + v h(x_i))] \\ &= \sum_{i=1}^n w_i^m \exp[-y_i v h(x_i)] & w_i^m \stackrel{\text{def}}{=} \exp[-y_i f_{m-1}(x_i)] \\ &= \sum_{i=1}^n w_i^m [\mathbb{I}(y_i = h(x_i)) e^{-v} + \mathbb{I}(y_i \neq h(x_i)) e^v] & h(x_i) \in \{1, -1\} \\ &= \sum_{i=1}^n w_i^m [(e^v - e^{-v}) \mathbb{I}(y_i \neq h(x_i)) + e^{-v}] & \mathbb{I}(y_i = h(x_i)) = 1 - \mathbb{I}(y_i \neq h(x_i)) \end{aligned}$$

- so i mean the above is mainly algebra i am not sure if there is that much to it except for a helpful way to re-write the objective
- there is a lot of kind of messy algebra in this section that i am not sure is super useful
- the real take away is basically this a generalization of ada boost
- in practice this has a robustness issue since ada boost is not robust to outliers

review

- so far we have seen that using a basis function to obtain nonlinear model $f(x) = \sum_{m=1} v_m h_m(x)$ if you know the basis functions
- could use adaptive basis function models if you do not know the basis.
- and forward stage wise additive modeling greedily fits h_m to minimize average loss
- but fsam only works for some loss functions
- we need a more general model

gradient boosting / any boost

FSAM with squared loss

- our objective function is

$$j(v, h) = \frac{1}{n} \sum_{i=1}^n (y_i - (f_{m-1}(x_i) + v h(x_i)))^2$$

if \mathcal{H} is closed under scaling then we can just set $v = 1$ and maximize and the model will adjust its self

- doing $v=1$ yields

$$j(h) = \frac{1}{n} \sum_i ([y_i - f_{m-1}(x_i) - h(x)]^2$$

this is equivalent to fitting function residuals with least squares regression

- so we can think of this as sequentially building models such that each one minimizes least squared residuals with the ones before set

interpret the residual

- so our objective $J(f) = \frac{1}{n} \sum_{i=1} (y_i - f(x_i))^2 = \frac{1}{n} \sum_{i=1} (y_i - f_{m-1}(x) - f_m(x_i))^2$
- so we can see that $\frac{\partial J(j)}{\partial f(x_i)} = -2(y_i - f(x_i))$ this gradient with respect to f is saying how should we try to change the output of f to minimize square loss
- so in other words our residual is the gradient
- so what we are doing at each step is learning $h \in \mathcal{H}$ to fit the residual

$$f \leftarrow f + v h$$

functional gradient descent

- we want to minimize our objective

$$J(f) = \sum_{i=1}^n \ell(y_i, f(x_i))$$

- not that $j(f)$ only depends at f evaluated at n training points that is $f = (f(x_1) \cdots f(x_n))^t$ so treating these as parameters we can write

$$j(f) = \sum \ell(y_i, f_i)$$

- the negative gradient $-g$ is the vector of partial derivatives of the j with respect to $f_i = f(x_i)$
- with gradient descent the final predictor will be

$$f_0 + \sum v_t(-g_t)$$

that is at every step we are updating our parameters in the direction of the gradient

- the unconstrained step direction $-g$ is called the “pseudo residual”
- so we only have n points which to use to estimate $h \in \mathcal{H}$ so we do projected least squares regression where

$$\min_{h \in \mathcal{H}} \sum (-g_i - h(x_i))^2$$

recap

- so we have the following objects

1. our objective function

$$j(f) = \sum \ell(y_i, f(x_i))$$

2. unconstrained gradient $g \in \mathbb{R}^n$ wrt $f = (f(x_1) \cdots f(x_n))^t$ (so this is a true gradient of the loss function with respect to what we have learned so far)
3. then we have the projected negative gradient $h \in \mathcal{H}$ so that is

$$h = \operatorname{argmin}_{h \in \mathcal{H}} \sum_{i=1}^n (-g_i - h(x_i))^2$$

4. and we update our function at each step according to this projected gradient $f \leftarrow f + vh$

Gradient boosting algorithm

1 Initialize f to a constant: $f_0(x) = \arg \min_{\gamma} \sum_{i=1}^n \ell(y_i, \gamma)$.

2 For m from 1 to M :

3 Compute the pseudo-residuals (negative gradient):

$$r_{im} = - \left[\frac{\partial}{\partial f(x_i)} \ell(y_i, f(x_i)) \right]_{f(x_i) = f_{m-1}(x_i)} \quad (26)$$

4 Fit a base learner h_m with squared loss using the dataset $\{(x_i, r_{im})\}_{i=1}^n$.

5 [Optional] Find the best step size $v_m = \arg \min_v \sum_{i=1}^n \ell(y_i, f_{m-1}(x_i) + v h_m(x_i))$.

6 Update $f_m = f_{m-1} + \lambda v_m h_m$

7 Return $f_M(x)$.

binomial boost with logistic loss

- recall that logistic loss with $y \in [-1, 1]$

$$\ell(y, f(x)) = \log(1 + e^{-y f(x)})$$